

# Income Case Study

# Data Challenges

# Missing Values


For numeric fields (age, education-num), we used **median imputation** and added **missing-value indicator flags** to preserve information.

For education\_level, missing entries were **replaced** with "Unknown" to avoid losing categorical information.

```
# Handle missing feature values
if "age" in df.columns:
    df["age_missing_flag"] = df["age"].isna().astype("int8")
    df["age"] = df["age"].fillna(df["age"].median()).astype("int64")

# education_num: median imputation + int
if "education-num" in df.columns:
    df["edu_num_missing_flag"] = df["education-num"].isna().astype("int8")
    df["education-num"] = (
        df["education-num"]
            .fillna(df["education-num"].median())
            .astype("int64")
    )

# education_level: fill missing with 'Unknown'
if "education_level" in df.columns:
    df["education_level"] = df["education_level"].fillna("Unknown")
```



# Imbalanced Classes

- **Catboost cost sensitive learning**


```
pos_weight = float((len(y_train) -  
y_train.sum()) / y_train.sum())
```

```
cat_model = CatBoostClassifier(...  
class_weights=[1.0, pos_weight],  
...)
```

- **Evaluation metrics that make sense with imbalance data:**

ROC AUC, PR AUC, Precision & Recall at the %9.5 cutoff


Next Steps: Try calibrated classifier to adjust predicted probabilities



# Wealth Bias in the Labeled Sample

**“We expect total of 2 million customers that fulfills the target out of the banks 21 million total customer base”**

Our labeled sample is much richer than the real bank, so raw probabilities are too optimistic.



# Name

Their high-income rates are extreme:

- Bernie: ~97% high income
- Johanna: ~97%
- Audrey, Billie, Karen: 0% high income
- Also, missing income is mostly on Audrey, Billie, Karen; almost never on Bernie/Johanna.


Name encodes the target by construction (**data leak**). If we keep it, the model will just learn ‘Bernie/Johanna = rich’ and won’t generalise. Therefore, we drop Name entirely before modeling.

# Feature Engineering

# Capital Gain & Loss

These are very skewed: most people have 0, but when non-zero, it's a huge signal.

- **capital-gain > 0**
  - ~9% of people
  - high income rate ≈ 66.8%
  - about 21% of all >50K customers.
- **capital-loss > 0**
  - ~5% of people
  - high income rate ≈ 55.7%



Any sign of investment gains/losses is a strong indicator of wealth, but only for a minority of customers.



```
# ----- capital features -----  
df_model["has_capital_gain"] = (df_model["capital-gain"] > 0).astype(int)  
df_model["has_capital_loss"] = (df_model["capital-loss"] > 0).astype(int)
```

# Binned features

Captures meaningful ranges instead of treating every number separately.

```
# ----- age features -----
age_bins    = [0, 25, 35, 45, 55, 65, 120]
age_labels  = ["<25", "25-34", "35-44", "45-54", "55-64", "65+"]
df_model["age_group"] = pd.cut(df_model["age"], bins=age_bins, labels=age_labels, right=False)


# ----- hours features -----
hours_bins   = [0, 35, 41, 46, 61, 200]
hours_labels = ["<35", "35-40", "41-45", "46-60", "60+"]
df_model["hours_group"] = pd.cut(df_model["hours-per-week"], bins=hours_bins,
                                  labels=hours_labels, right=False)
df_model["long_hours_flag"] = (df_model["hours-per-week"] >= 46).astype(int)
```



# Education score

to express a **strong monotonic relationship**

We convert ordered categories into a single numeric score so the model can learn this trend directly. This reduces noise and improves consistency across similar education levels.




```
# ----- education features -----
edu_order = [
    "Preschool", "1st-4th", "5th-6th", "7th-8th", "9th", "10th", "11th", "12th",
    "HS-grad", "Some-college", "Assoc-voc", "Assoc-acdm",
    "Bachelors", "Masters", "Prof-school", "Doctorate"
]
edu_map = {edu: i for i, edu in enumerate(edu_order)}
df_model["education_level"] = df_model["education_level"].fillna("Unknown")
df_model["education_score"] = df_model["education_level"].map(edu_map).fillna(-1)
```

# Interaction Effects #1

## Marital Status + Age + Hours Worked

- Married + Age 40–60 + 45+ hrs/week → >60% high-income rate
- Never-married + Age <35 → <10%

```
married_statuses = {"Married-civ-spouse", "Married-AF-spouse", "Married-spouse-absent"}  
  
# High-income profile: Married + Age 40–60 + 45+ hours/week  
df_model["married_midage_longhours"] = (  
    df_model["marital-status"].isin(married_statuses)  
    & (df_model["age"] >= 40)  
    & (df_model["age"] <= 60)  
    & (df_model["hours-per-week"] >= 45)  
).astype(int)  
  
# Low-income profile: Never-married + Age <35  
df_model["young_never_married"] = (  
    (df_model["marital-status"] == "Never-married")  
    & (df_model["age"] < 35)  
).astype(int)
```




# Interaction Effects #2

## Relationship + Sex

- “Husband” relationship code → ~45–55%
- “Own-child” → <5%


```
# High-income household heads
df_model["rel_husband_male"] = (
    (df_model["relationship"] == "Husband") &
    (df_model["sex"] == "Male")
).astype(int)

df_model["rel_wife_female"] = (
    (df_model["relationship"] == "Wife") &
    (df_model["sex"] == "Female")
).astype(int)
```



# Marital Status

To express similarity add grouped features of “is\_married”, “is\_in\_couple” and “has\_children”.




```
# ----- marital / relationship -----
married_statuses = {"Married-civ-spouse", "Married-AF-spouse", "Married-spouse-absent"}
df_model["is_married"] = df_model["marital-status"].isin(married_statuses).astype(int)
df_model["is_in_couple"] = df_model["relationship"].isin({"Husband", "Wife"}).astype(int)
df_model["has_children"] = df_model["relationship"].isin({"Own-child", "Other-relative"}).astype(int)
```

# Is American

The native-country distribution is extremely imbalanced toward the U.S.

```
# ----- geography -----
df_model[ "is_US" ] = (df_model[ "native-country" ] == "United-States").astype(int)
```




# Region

Reduces Noise from Rare Categories

Income Patterns Differ Strongly by Region.

Region feature helps capture macro-economic effects that single-country categories miss.



```
# Simple region grouping example
europe = {"England", "Germany", "Portugal", "France", "Italy", "Ireland", "Scotland", "Greece", "Holland-Netherlands"}
latin_america = {"Mexico", "Cuba", "Jamaica", "Honduras", "Puerto-Rico", "Columbia", "Ecuador",
                 "Trinidad&Tobago", "Peru", "Nicaragua", "El-Salvador", "Guatemala", "Haiti", "Dominican-Republic"}
asia = {"Philippines", "India", "Japan", "China", "Vietnam", "Laos", "Cambodia", "Thailand", "Taiwan", "Hong Kong"}

def map_region(c):
    if c in europe: return "Europe"
    if c in latin_america: return "LatAm"
    if c in asia: return "Asia"
    if c == "United-States": return "US"
    return "Other"

df_model["country_region"] = df_model["native-country"].map(map_region)
```

# Model

# Why Catboost?

Easy to use with categorical data

Class-weights feature

“PROUC” performance measure

# Hyperparameter Optimization

## CV + Grid Search + Optuna

**Search space design:** We exposed the usual high-leverage CatBoost knobs: depth, learning\_rate, iterations, regularization (l2\_leaf\_reg), and maybe some tree-level constraints / sampling rates. We kept the space reasonably tight for Optuna, based on prior experience and a few manual runs, so we're not wasting budget in obviously bad regions.

**Stage 1 – Grid Search:** provides sweet starting point for Optuna

**Stage 2 – Optuna (Bayesian / TPE + pruning):** Provides even better hyperparameters

```
Best params : {'depth': 6, 'learning_rate': 0.06284728302324294, 'l2_leaf_reg': 2.3532353602061975,  
'iterations': 1179}
```

# How we evaluated the model?

Goal: “Predict which customers have annual income > \$50k so we can select ~2M of 21M customers (~9.5%) for a premium campaign.”

When we pick the top 10% highest-scoring customers, 96% of them are truly high-income.

Top-K precision / recall (how well we find high-income customers):

Top 1% -> precision=1.0000, recall=0.0352

Top 5% -> precision=0.9975, recall=0.1762

Top 10% -> precision=0.9735, recall=0.3443

Top 20% -> precision=0.8449, recall=0.5977

VALIDATION

Top-K precision / recall (how well we find high-income customers):

Top 1% -> precision=1.0000, recall=0.0353



Top 5% -> precision=0.9924, recall=0.1754

Top 10% -> precision=0.9622, recall=0.3405


Top 20% -> precision=0.8588, recall=0.6078

TESTING


# Metrics #1: ROC-AUC, Precision-Recall AUC



# Metrics #2: Confusion Matrix, Precision, Recall, F1



Accuracy	: 0.8403
Precision	: 0.6708
Recall	: 0.8541
F1-score	: 0.7514




Accuracy	: 0.8349
Precision	: 0.6579
Recall	: 0.8657
F1-score	: 0.7476

# Wealth Overestimation in Labeled sample

Among the customers where we know income, about 28% are above \$50k. But you told us that in the full bank, only about 9.5% are above \$50k.

So if we just trained a model on the labeled sample, it would think the world is richer than it really is and would assign probabilities that are too high. We keep the ranking — who looks richer than whom — but we shrink the probabilities so that, on average, they match the 9.5% reality. That's what calibration does.



# Pick Customers with the best probability

“We expect total of 2 million customers that fulfills the target out of the banks 21 million total customer base”

%9.5 percent of the customers with the best probability of being high income should be chosen when model is applied to main dataset.

# What to improve

Better handling of missing values (e.g. KNN imputer)

Trying other algorithms and comparing metrics

Simulate the real 9.5% positive rate by downsampling in testing part

Simplifying features for easier deployment

Voting Classifier or Ensemble Model